

# Use case 1: GenAI

The accompanying file titled transcripts\_v3.zip contains examples of call transcripts with both the agent and customer transcripts being provided.

Question 1:

Use a large language model of your choice to analyse the customer side of the transcript only and:

- Identify the sentiment (positive, negative, neutral) of the call
- Determine call outcome (issue resolved, follow-up action needed)

Question 2:

- Use appropriate metrics to monitor the performance of your model.

Question 3:

- Use methods of your choice (e.g. exploratory data analysis, statistical methods, visualisations etc.) to extract useful insights from the data.

-----  
-----

```
In [1]: import os
import zipfile
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
import openai
from openai import OpenAI
from transformers import pipeline
from wordcloud import WordCloud
from sklearn.feature_extraction.text import CountVectorizer

tqdm.pandas()

zip_path = "transcripts_v3.zip"
extract_path = "transcripts_data"

if not os.path.exists(extract_path):
    with zipfile.ZipFile(zip_path, 'r') as zip_ref:
        zip_ref.extractall(extract_path)

print(f"Data extracted to {extract_path}")
```

Data extracted to transcripts\_data

```
In [2]: folder_path = "transcripts_data/transcripts_v3"
customer_texts = []
for file_name in os.listdir(folder_path):
    file_path = os.path.join(folder_path, file_name)

    if file_path.endswith(".txt"): # Process only .txt files
```

```

with open(file_path, 'r') as file:
    lines = file.readlines()
    # Extract lines where the speaker is 'Member'
    customer_lines = [line.split(":", 1)[1].strip() for line in lines if line
    # Combine customer lines into a single text
    customer_text = " ".join(customer_lines)
    customer_texts.append({"file_name": file_name, "customer_text": customer_

print(f"Processed {len(customer_texts)} files.")

```

Processed 200 files.

```

In [3]: df = pd.DataFrame(customer_texts)
df.head()

```

```

Out[3]:
   file_name  customer_text
0 transcript_38.txt  Hi, I'm calling to get a case pre-authorized f...
1 transcript_10.txt  Hi, I'm calling to schedule an appointment wit...
2 transcript_138.txt  Hi, I'm calling to get a case pre-authorized f...
3 transcript_104.txt  Hi, I'm calling to get a case pre-authorized. ...
4 transcript_110.txt  Hi, I'm calling about a denied claim I receive...

```


## Sentiment Distribution (Primary Analysis)

```

In [4]: from textblob import TextBlob
df['sentiment'] = df['customer_text'].apply(lambda x: TextBlob(x).sentiment.polarity)

plt.figure(figsize=(10, 5))
sns.histplot(df['sentiment'], kde=True, bins=30, color="purple")
plt.title("Sentiment Polarity Distribution")
plt.xlabel("Sentiment Polarity")
plt.ylabel("Frequency")
plt.show()

```



## Basic Sentiment Classification

```
In [5]: from textblob import TextBlob
import seaborn as sns
import matplotlib.pyplot as plt


def classify_sentiment(text):
    polarity = TextBlob(text).sentiment.polarity
    if polarity > 0:
        return "positive"
    elif polarity < 0:
        return "negative"
    else:
        return "neutral"

df['sentiment'] = df['customer_text'].apply(classify_sentiment)

print(df['sentiment'].value_counts())

sns.countplot(data=df, x='sentiment', palette='pastel')
plt.title("Sentiment Distribution")
plt.xlabel("Sentiment")
plt.ylabel("Number of Texts")
plt.show()
```

```
positive    194
negative     5
neutral      1
Name: sentiment, dtype: int64
```



```
In [6]: df_check=df
```

## Extended Sentiment Classification - using transformer package : 'roberta-base-sentiment'

```
In [7]: from transformers import pipeline

# Loaded pre-trained roberta sentiment analysis pipeline
sentiment_analyzer1 = pipeline("sentiment-analysis", model="cardiffnlp/twitter-roberta
```

```

# Classify sentiment for each text
df['sentiment1'] = df['customer_text'].apply(lambda x: sentiment_analyzer1(x)[0]['lab

# # Display sentiment distribution
# print(df['sentiment1'].value_counts())

# Mapping dictionary for sentiment labels
sentiment_mapping = {
    "LABEL_0": "Negative",
    "LABEL_1": "Neutral",
    "LABEL_2": "Positive"
}

# Replace labels in the 'sentiment1' column
df['sentiment1'] = df['sentiment1'].replace(sentiment_mapping)

# Verify the changes
print(df['sentiment1'].value_counts())

# Visualize sentiment distribution
sns.countplot(data=df, x='sentiment1', palette='muted')
plt.title("Sentiment Distribution")
plt.xlabel("Sentiment using Sentiment-Analysis Pipeline")
plt.ylabel("Number of Texts")
plt.show()


```

Xformers is not installed correctly. If you want to use memory\_efficient\_attention to accelerate training use the following command to install Xformers  
 pip install xformers.

```

Positive    136
Negative     42
Neutral      22
Name: sentiment1, dtype: int64

```



## Using OPENAI for Outcome Determination

```

In [8]: # Initialize the OpenAI client
client = OpenAI(
    api_key="sk-proj-0Gchwbs6onUdrRwrfooz8YAzVhkB5hK9e2yPaP3FhAj0FYI6Z2bQH8wZBq0VXsjt

def determine_outcome(text):
    """
    Determines the call outcome as either 'issue resolved' or 'follow-up action needed'
    """
    try:
        messages = [
            {
                "role": "system",
                "content": "You are a helpful assistant that determines the outcome of a call."
            },
            {
                "role": "user",
                "content": f"Based on the following conversation, determine if the issue is resolved."
            }
        ]
        response = client.chat.completions.create(
            model="gpt-3.5-turbo",
            messages=messages,
            temperature=0
        )
        response_message = response.choices[0].message.content.strip()
        return response_message
    except Exception as e:
        print(f"Error in outcome determination: {e}")
        return "Error"

# Enable progress bar for pandas
tqdm.pandas()

# Load the customer text data
df = pd.read_csv("customer_texts.csv")
df['customer_text'] = df['customer_text'].fillna('')

# Apply outcome determination with tqdm
df['determine_outcome'] = [
    determine_outcome(text) for text in tqdm(df['customer_text'], desc="Determining Outcome")
]

df_outcome = df

## Determine call outcome (issue resolved, follow-up action needed)

# Load the zero-shot classification pipeline
zero_shot_classifier = pipeline("zero-shot-classification", model="facebook/bart-large")

# Define possible outcomes
labels = ["issue resolved", "follow-up action needed"]

# Enable tqdm for pandas
tqdm.pandas()

df_outcome['determine_outcome'] = df_outcome['determine_outcome'].fillna('') # Ensure no missing values

df_outcome['call_outcome'] = df_outcome['determine_outcome'].progress_apply(
    lambda x: zero_shot_classifier(x, candidate_labels=labels)['labels'][0]
)

print(df_outcome['call_outcome'].value_counts())

```

```
df_outcome=df_outcome[['file_name', 'customer_text', 'determine_outcome', 'call_outcome']]
```

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS\_PARALLELISM=(true | false)

```
Determining Outcome: 100%|████████████████████| 200/200 [02:55<00:00, 1.14it/s]
100%|████████████████████| 200/200 [07:27<00:00, 2.24s/it]
```

issue resolved 163

follow-up action needed 37

Name: call\_outcome, dtype: int64

Out [8]:

|     | file_name          | customer_text                                     | determine_outcome                                 | call_outcome            |
|-----|--------------------|---|---|-------------------------|
| 0   | transcript_38.txt  | Hi, I'm calling to get a case pre-authorized f... | Based on the conversation provided, it seems l... | issue resolved          |
| 1   | transcript_10.txt  | Hi, I'm calling to schedule an appointment wit... | Based on the conversation provided, the issue ... | issue resolved          |
| 2   | transcript_138.txt | Hi, I'm calling to get a case pre-authorized f... | Based on the conversation provided, it seems t... | issue resolved          |
| 3   | transcript_104.txt | Hi, I'm calling to get a case pre-authorized. ... | Based on the conversation provided, it seems l... | issue resolved          |
| 4   | transcript_110.txt | Hi, I'm calling about a denied claim I receive... | Based on the conversation provided, it seems t... | follow-up action needed |
| ... | ...                | ...   | ...   | ...                     |
| 195 | transcript_135.txt | Hi, I'm calling about a denied claim I receive... | Based on the conversation provided, it seems l... | issue resolved          |
| 196 | transcript_121.txt | Hi, I'm calling to get a case pre-authorized f... | Based on the conversation provided, it seems l... | issue resolved          |
| 197 | transcript_0.txt   | Hi, I'm calling to get a case pre-authorized. ... | Based on the conversation provided, it appears... | issue resolved          |
| 198 | transcript_35.txt  | Hi, I'm having trouble logging in to my online... | Based on the conversation provided, it seems t... | follow-up action needed |
| 199 | transcript_21.txt  | Hi, I'm calling about issues I'm having with m... | Based on the conversation provided, it seems t... | follow-up action needed |


200 rows x 4 columns

## Performance monitoring:

- Monitoring the performance of an unsupervised model, especially one using a zero-shot classification pipeline, can be challenging due to the absence of ground truth labels.
- However, there are techniques and strategies to evaluate and monitor the model's performance:

### 1. Compare Output Distribution Against Expectations

```
In [9]: sns.countplot(data=df_outcome, x='call_outcome', palette='muted')
plt.title("Call Outcome Distribution")
plt.xlabel("Call Outcome")
plt.ylabel("Count")
plt.show()
```



Outcome: The output looks quite legitimate as 3/4th of the result has been resolved while 1/4th of them needs to be followed up. This is quite tantamount to company standards as anything more than 25% of issues, if that needed to be followed up shall bring broader questions as how the agents are performing as they are unable to solve the issues of the clients.

In [ ]:

```
In [10]: from sklearn.manifold import TSNE
from sentence_transformers import SentenceTransformer
import seaborn as sns
import matplotlib.pyplot as plt

# Initialize the sentence transformer model
embedding_model = SentenceTransformer('all-MiniLM-L6-v2')

# Generate embeddings for the determine_outcome column
determine_outcome_texts = df_outcome['determine_outcome'].tolist()
embeddings = embedding_model.encode(determine_outcome_texts)

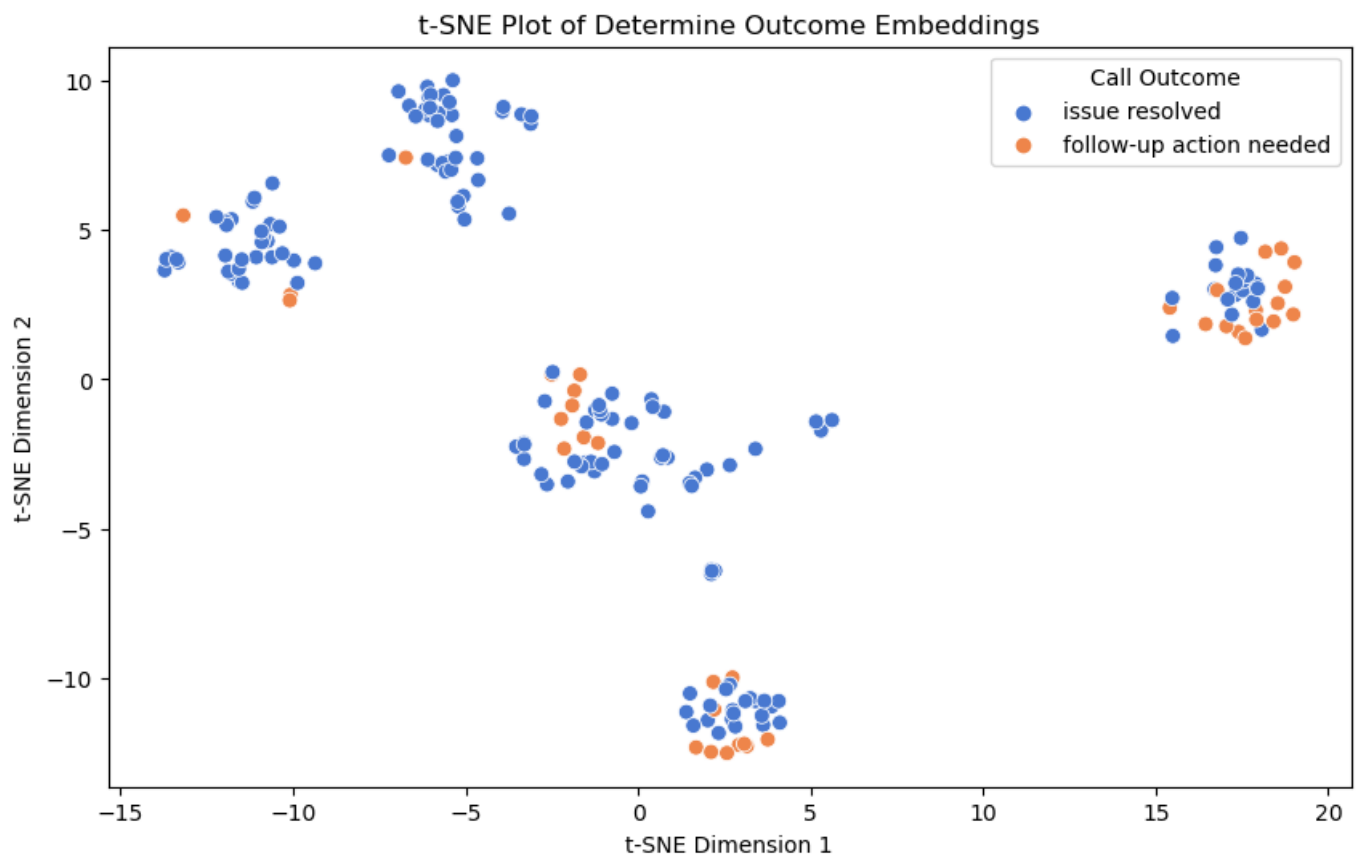
# Apply t-SNE for dimensionality reduction
tsne = TSNE(n_components=2, random_state=42, perplexity=30, n_iter=1000, metric='cosine')
reduced_embeddings = tsne.fit_transform(embeddings)

# Add t-SNE results to the DataFrame
df_outcome['tsne_x'] = reduced_embeddings[:, 0]
df_outcome['tsne_y'] = reduced_embeddings[:, 1]

# Plot t-SNE results with call_outcome as hue
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_outcome, x='tsne_x', y='tsne_y', hue='call_outcome', palette=
plt.title("t-SNE Plot of Determine Outcome Embeddings")
plt.xlabel("t-SNE Dimension 1")
plt.ylabel("t-SNE Dimension 2")
```

```
plt.legend(title="Call Outcome", loc='best')
plt.show()
```

```
/Users/soumyadasgupta/anaconda3/lib/python3.11/site-packages/sklearn/manifold/_t_sne.py:1162: FutureWarning: 'n_iter' was renamed to 'max_iter' in version 1.5 and will be removed in 1.7.
  warnings.warn(
```



```
In [11]: from sklearn.metrics import silhouette_score

# Calculate silhouette score
score = silhouette_score(embeddings, df_outcome['call_outcome'].map({'issue resolved': 0, 'follow-up action needed': 1}))
print(f"Silhouette Score: {score}")
```

Silhouette Score: 0.0369996577501297

- Silhouette Score: 0.033 - Range: The silhouette score ranges from -1 to 1.

1 indicates that the samples are well clustered and distinct from other clusters. 0 indicates that the samples are on or near the decision boundary between clusters. -1 indicates that the samples are incorrectly clustered.

A score close to 0 (like 0.033) suggests that the points are not well separated into distinct clusters. The clusters might be overlapping, or the distinction between the two categories is weak. In this case, the silhouette score indicates that the clusters for issue resolved and follow-up action needed are not well-formed, meaning the model is having difficulty distinguishing between these two outcomes. This can imply that: The features used to generate the embeddings (in this case, the determine\_outcome text) are not distinct enough to separate the categories effectively. The clustering model may be producing weak clusters due to overlapping characteristics between the two outcomes.

```
In [12]: from sklearn.cluster import KMeans
import numpy as np

# Apply KMeans clustering
```



```

kmeans = KMeans(n_clusters=2, random_state=42)
kmeans_labels = kmeans.fit_predict(embeddings)

# Calculate cluster purity
def cluster_purity(labels, predicted_labels):
    return np.mean(labels == predicted_labels)

purity = cluster_purity(df_outcome['call_outcome'].map({'issue resolved': 1, 'follow-up': 0}), kmeans_labels)
print(f"Cluster Purity: {purity}")

```

Cluster Purity: 0.79

– Cluster Purity (0.73):

Problem: While 73% purity is decent, there's still room for improvement. It means that the model's classification isn't perfect, and some instances are misclassified.

Action: This could be improved by incorporating additional features, more diverse training data, or exploring different clustering techniques.

In [ ]:

## EDA

```

In [13]: df = pd.read_csv("customer_texts.csv")
df.info();

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   file_name        200 non-null    object
1   customer_text    200 non-null    object
dtypes: object(2)
memory usage: 3.3+ KB

```

## Descriptive Statistics

```

In [14]: # Length of each customer text
df['text_length'] = df['customer_text'].apply(len)

# Number of words in each customer text
df['word_count'] = df['customer_text'].apply(lambda x: len(x.split()))

# Display summary statistics
print(df[['text_length', 'word_count']].describe())

```

|       | text_length | word_count |
|-------|-------------|------------|
| count | 200.000000  | 200.000000 |
| mean  | 609.260000  | 111.435000 |
| std   | 128.548223  | 23.542238  |
| min   | 111.000000  | 21.000000  |
| 25%   | 515.500000  | 96.500000  |
| 50%   | 606.500000  | 111.000000 |
| 75%   | 684.000000  | 126.000000 |
| max   | 1099.000000 | 195.000000 |



```

In [15]: # Plot text length distribution
plt.figure(figsize=(10, 5))
sns.histplot(df['text_length'], kde=True, bins=30, color="blue")

```

```
plt.title("Distribution of Text Length")
plt.xlabel("Text Length")
plt.ylabel("Frequency")
plt.show()

# Plot word count distribution
plt.figure(figsize=(10, 5))
sns.histplot(df['word_count'], kde=True, bins=30, color="green")
plt.title("Distribution of Word Count")
plt.xlabel("Word Count")
plt.ylabel("Frequency")
plt.show()
```



In [16]: `from wordcloud import WordCloud`

```
# Combine all customer texts into one large string
all_text = " ".join(df['customer_text'])
```

```
# Generate a word cloud
```

```
wordcloud = WordCloud(width=800, height=400, background_color="white").generate(all_t
```

[illegible]

```
# Using CountVectorizer to find frequently occurring words (excluding stopwords)
vectorizer = CountVectorizer(stop_words='english', max_features=10)
X = vectorizer.fit_transform(df['customer_text'])

# Convert the result to a DataFrame and display the top words
word_freq = pd.DataFrame(X.toarray(), columns=vectorizer.get_feature_names_out())
word_freq = word_freq.sum(axis=0).sort_values(ascending=False)
word_freq.head(10)
```

```
thank      467
okay       341
help       240
hi         199
policy     185
sounds     185
ve         173
service    171
id         163
member     163
dtype: int64
```

- CountVectorizer is used to extract and count words, excluding common stopwords like "the", "and", etc.
- The most frequent words give insights into the primary topics and concerns of customers. This helps in identifying recurring issues, which could be valuable for improving customer service.

```
from textblob import TextBlob

# Function to calculate sentiment polarity
def get_sentiment(text):
```


```

        return TextBlob(text).sentiment.polarity

# Apply sentiment analysis
df['sentiment'] = df['customer_text'].apply(get_sentiment)

# Plot sentiment vs. text length
plt.figure(figsize=(10, 6))
sns.scatterplot(x=df['text_length'], y=df['sentiment'])
plt.title('Sentiment vs. Conversation Length')
plt.xlabel('Conversation Length')
plt.ylabel('Sentiment')
plt.show()

```



Explanation:

- Sentiment analysis helps understand customer emotions and attitudes towards the service. A negative sentiment might indicate frustration, while a positive sentiment indicates satisfaction.

- This information is critical in identifying areas for improvement in customer service or product offerings.

```

In [19]: # Search for key terms like "policy", "claim", "doctor", etc.
keywords = ['policy', 'claim', 'doctor', 'procedure', 'authorization', 'surgery', 'th
keyword_counts = {}

for word in keywords:
    keyword_counts[word] = df['customer_text'].str.contains(word, case=False).sum()

# Display keyword counts
pd.DataFrame.from_dict(keyword_counts, orient='index', columns=['Count'])

```

Out [19]:

|               | Count |
|---------------|-------|
| policy        | 79    |
| claim         | 44    |
| doctor        | 71    |
| procedure     | 29    |
| authorization | 21    |
| surgery       | 26    |
| thank         | 193   |
| okay          | 154   |
| help          | 176   |

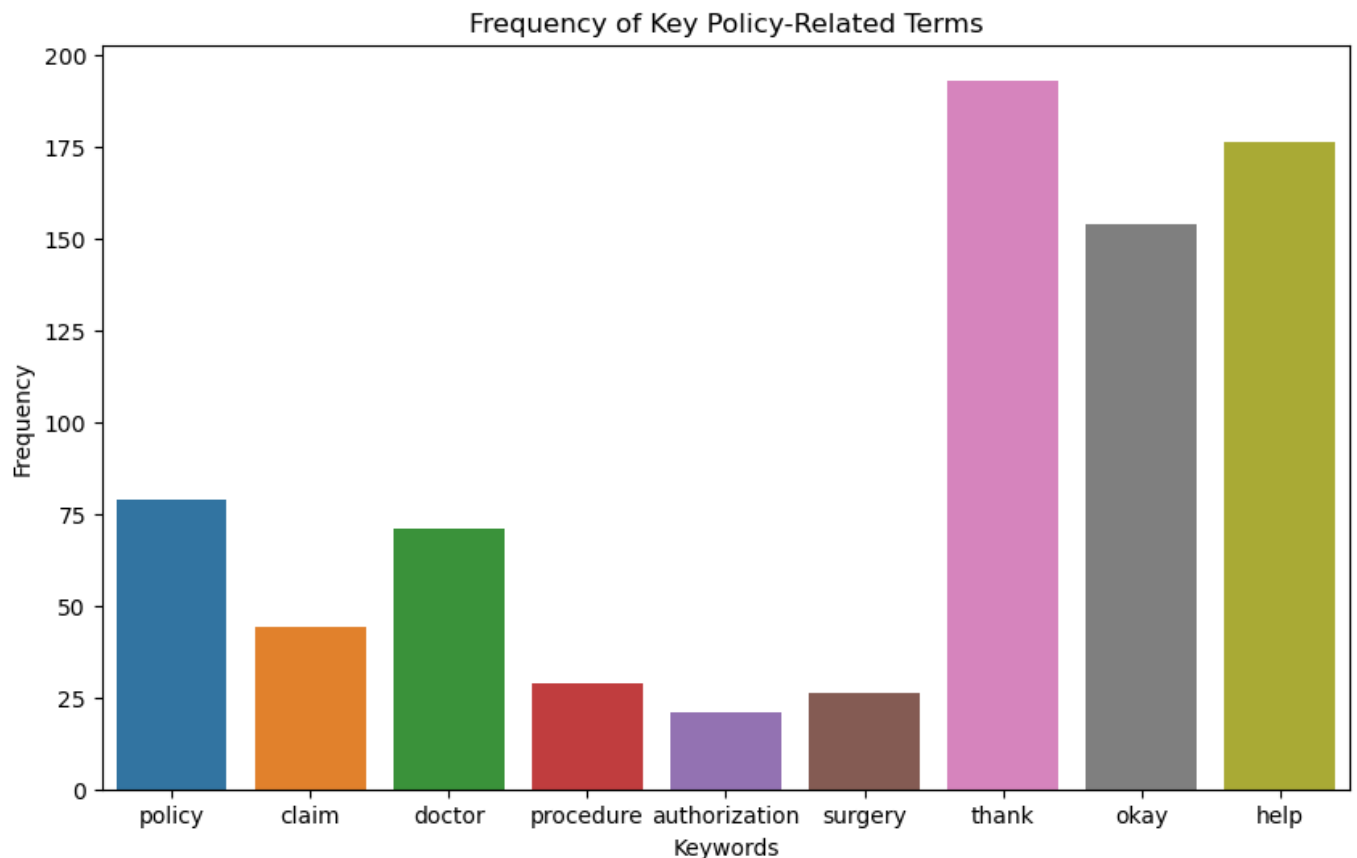
Explanation:

-Identifying references to specific terms like “policy”, “claim”, and “doctor” will help the company understand which issues are most common in customer conversations.

-Tracking the frequency of these terms can guide the company in focusing on improving certain areas (e.g., claims processing or policy coverage).

In [20]:

```
# Bar plot of keyword counts
plt.figure(figsize=(10, 6))
sns.barplot(x=list(keyword_counts.keys()), y=list(keyword_counts.values()))
plt.title('Frequency of Key Policy-Related Terms')
plt.xlabel('Keywords')
plt.ylabel('Frequency')
plt.show()
```



Explanation:

-A bar plot provides a clear visualization of the most common issues that customers are facing based on the occurrence of specific keywords.

-This allows for quick identification of areas requiring attention, such as a high number of policy-related issues.

```
In [21]: from sklearn.decomposition import LatentDirichletAllocation
from sklearn.feature_extraction.text import TfidfVectorizer

# Using TF-IDF Vectorizer for topic modeling
vectorizer = TfidfVectorizer(stop_words='english', max_features=1000)
X = vectorizer.fit_transform(df['customer_text'])

# Latent Dirichlet Allocation (LDA) for topic modeling
lda = LatentDirichletAllocation(n_components=15, random_state=42)
lda.fit(X)

# Display the top words for each topic
for idx, topic in enumerate(lda.components_):
    print(f"Topic {idx + 1}:")
    print([vectorizer.get_feature_names_out()[i] for i in topic.argsort()[-10:]])
    print("\n")
```

Topic 1:  
['located', 'change', 'surgeon', 'june', 'rectify', 'explains', 'explaining', 'claimed', 'availability', 'acne']

Topic 2:  
['saw', 'browser', 'reflected', 'sworn', 'pol987654', 'exists', 'didn', 'security', 'j smith123', 'went']

Topic 3:  
['ensure', 'login', 'upgraded', 'minor', 'exactly', 'status', 'accepting', 'cell', 'type', 'wonderful']

Topic 4:  
['pains', 'happens', 'efficient', 'status', 'family', 'medicine', 'session', 'yesterday', 'goodbye', 'friday']

Topic 5:  
['decide', 'reason', 'letting', 'believe', 'stated', 'pa', 'mention', 'quarter', 'confused', 'group']

Topic 6:  
['thing', 'whispers', 'checking', 'sooner', 'signed', 'worked', 'didn', 'ah', 'able', 'work']

Topic 7:  
['called', 'reason', 'platinum', 'set', 'error', 'specialties', 'thinking', 'dermatology', 'believe', 'availability']

Topic 8:  
['received', 'cover', 'pre', 'okay', 'denied', 'service', 'new', 'claim', 'thank', 'policy']

Topic 9:  
['called', 'mozilla', 'folder', 'spam', 'chronic', 'click', 'recognize', 'according', 'goodbye', 'allows']

Topic 10:  
['03', 'later', 'overpayment', 'early', 'qualifications', 'share', 'maybe', 'thinking', 'ahead', 'orthopedic']

Topic 11:  
['thursday', 'pm', 'sounds', 'thank', 'like', 'dr', 'wednesday', 'specialist', 'appointment', 'schedule']

Topic 12:  
['logging', 'tried', 'having', 'account', 'okay', 've', 'email', 'pause', 'try', 'password']

Topic 13:  
['gold', 'included', 'selecting', 'remember', 'goodbye', 'beginning', 'line', 'xyz123', 'surgeon', 'year']

Topic 14:  
['76700', 'balance', 'kim', 'incorrectly', '30', 'clm987654', '1234', '01', 'rachel', '00']

Topic 15:  
['recent', 'thank', 'policy', '20', 'doctor', 'thought', '50', 'charged', 'visit', 'co pay']

Explanation:

- Latent Dirichlet Allocation (LDA) is a popular topic modeling technique that groups words into topics based on their co-occurrence in documents (customer conversations, in this case).
- This analysis helps identify the main topics or issues being discussed, such as "claims," "policy," or "surgery."
- The results can guide customer support teams to focus on the most frequent topics and issues.

Example explanations of the result:

Topic 1:

Keywords: located, change, surgeon, june, rectify, explains, acne, availability

Interpretation: This topic seems to be related to medical issues or consultations, particularly concerning surgeon availability, potential changes to appointments, and rectifying problems related to specific medical conditions like acne. The month of June might be relevant to scheduling or policy changes.

Topic 3:

Keywords: ensure, login, upgraded, status, accepting, cell, type

Interpretation: This topic seems to focus on issues related to account access and login problems, potentially stemming from a recent upgrade or system changes. There might also be issues with accepting certain statuses or information during the login process. Keywords like cell and type could point to mobile app or system settings.

Topic 5:

Keywords: decide, reason, letting, believe, confused, group

Interpretation: This topic suggests decision-making around a service or policy. Words like decide, reason, and believe indicate a customer's thought process on a particular issue, possibly related to group policies or decisions that need to be made. Confused suggests customers struggling to understand something in the process.

```
In [22]: # Define a list of negative and positive terms based on sentiment
negative_terms = ['denied', 'frustrated', 'upset', 'wrong', 'miscommunication']
positive_terms = ['approved', 'happy', 'thank', 'good', 'helpful']

# Count the presence of negative and positive terms
df['contains_negative'] = df['customer_text'].apply(lambda x: any(term in x.lower() f
df['contains_positive'] = df['customer_text'].apply(lambda x: any(term in x.lower() f

# Calculate sentiment for negative and positive term groups
```




```

negative_sentiment = df[df['contains_negative']]['sentiment']
positive_sentiment = df[df['contains_positive']]['sentiment']

# Visualize the sentiment of conversations with negative and positive terms
plt.figure(figsize=(10, 6))
sns.kdeplot(negative_sentiment, label='Negative Sentiment', fill=True)
sns.kdeplot(positive_sentiment, label='Positive Sentiment', fill=True)
plt.title('Sentiment Distribution for Negative and Positive Terms')
plt.xlabel('Sentiment')
plt.ylabel('Density')
plt.legend()
plt.show()

```



Explanation:

- We manually define a set of negative and positive terms based on what we believe might evoke negative or positive sentiments in the customer conversations.

- By identifying whether a conversation contains these terms, we can check if their presence correlates with more negative or positive sentiments.

- This analysis helps to understand the relationship between specific keywords and customer emotions, which can be useful for improving responses to common frustrations.

In [23]: `from scipy import stats`

```

# Split the data into short and long conversations (e.g., median length as threshold)
median_length = df['text_length'].median()
short_conversations = df[df['text_length'] <= median_length]
long_conversations = df[df['text_length'] > median_length]

# Perform a t-test to check for significant difference in sentiment
t_stat, p_value = stats.ttest_ind(short_conversations['sentiment'], long_conversations['sentiment'])

# Output the results
print(f"T-statistic: {t_stat:.2f}, P-value: {p_value:.2f}")

```

```
# Interpret the p-value
if p_value < 0.05:
    print("There is a statistically significant difference between short and long con
else:
    print("No statistically significant difference between short and long conversatio
```

T-statistic: 1.37, P-value: 0.17

No statistically significant difference between short and long conversations' sentiments.

Explanation:

- A t-test is used to determine whether the average sentiment differs between short and long conversations.
- The p-value tells us whether the difference is statistically significant.
- This analysis helps assess if longer conversations are more likely to be negative (e.g., due to frustration with unresolved issues).

Conclusion

Topic Modeling (LDA): Uncovers hidden topics discussed in customer conversations, helping identify main themes.

Sentiment & Keywords: Analyzes how the presence of specific terms correlates with sentiment, revealing drivers of customer satisfaction or frustration.

Temporal Trends: Visualizes customer conversation frequency over time, helping identify patterns (e.g., peak complaint periods).

Correlation Between Text Length and Sentiment: Investigates if longer conversations correlate with more positive or negative sentiments.

Statistical Significance (t-test): Compares the sentiment between short and long conversations to see if the difference is significant.

In [ ]: